

# SQL Recap and Introduction to PostgreSQL

---

Sebastian Ernst, PhD


Course: Databases II, EAlilBISIS.li8K.5dfa09851a120.22

# Introduction

---

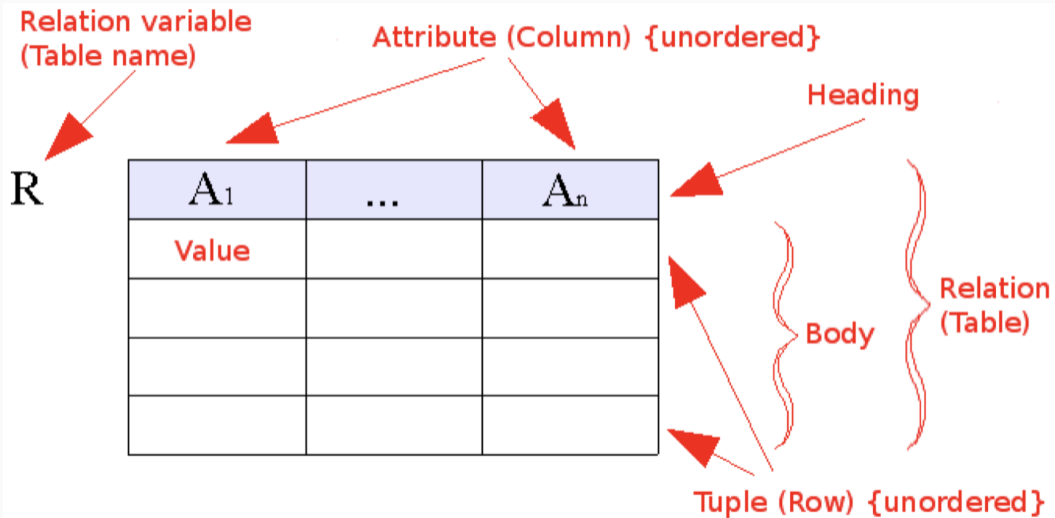
# The relational data model

- Proposed in 1969 by Edgar F. Codd.
- Models data as a collection of predicates over a finite set of variables.
- The conceptual model includes entities, their attributes and relationships between them (one-to-one, one-to-many, many-to-many).



Source: Wikipedia

# The relational data model



# The SQL Language

---

- SQL stands for Structured Query Language.
- Pronounced as 'ess-que-ell' or 'sequel'.
- Performs many roles:
  - data modelling language (DML),
  - data definition language (DDL),
  - data query language (DQL).

## SQL: creating relations

Relations are created using the CREATE TABLE function:

```
CREATE TABLE weather (  
    city      varchar(80),  
    temp_lo   int,           -- low temperature  
    temp_hi   int,           -- high temperature  
    prcp      real,         -- precipitation  
    date      date  
);
```

To delete a table, use DROP TABLE:

```
DROP TABLE weather;
```

## SQL: inserting data

Data is inserted using the INSERT statement:

```
INSERT INTO weather VALUES ('San Francisco', 46, 50, 0.25,  
    '1994-11-27');
```

If the set or order of columns is different than the table structure, they need to be specified in the query:

```
INSERT INTO weather (city, temp_lo, temp_hi, prcp, date)  
    VALUES ('San Francisco', 43, 57, 0.0, '1994-11-29');
```



**Selection** and **projection** are two fundamental operations in SQL.

**Selection** chooses which columns are included in the result set of a SELECT query:

```
SELECT surname, age FROM employees;
```

```
SELECT * FROM employees;
```

**Projection** chooses which rows are returned by means of the WHERE clause:

```
SELECT * FROM employees WHERE age > 30;
```

## SQL: aggregate functions

**Aggregate functions** combine a set of values into a single one. They include functions such as SUM, AVG, MIN, MAX, COUNT, etc.

```
SELECT AVG(age) FROM employees;
```

Aggregates may be computed for groups of rows instead of all records; these groups need to be explicitly defined.

```
SELECT AVG(age), department FROM employees GROUP BY department;
```

## SQL: filtering in groups

The WHERE clause always applies to individual records (i.e. prior to aggregation). To perform selection on *groups*, one uses the HAVING clause.

```
SELECT AVG(age), department FROM employees GROUP BY department
    WHERE age > 30;      -- use only persons older than 30 for averages
```

```
SELECT AVG(age), department FROM employees GROUP BY department
    HAVING AVG(age) > 30; -- only show depts where average is over 30
```

## SQL: inner joins

Joins merge **two sets of records**, creating a single result set.

Inner joins include only records which match both source sets:

```
SELECT * FROM weather, cities WHERE weather.city = cities.name;
```

```
SELECT * FROM weather INNER JOIN cities ON (weather.city = cities.name);
```

```
SELECT * FROM weather JOIN cities ON (weather.city = cities.name);
```

## SQL: outer joins

Outer joins take all records from the left, right or both source sets:

```
SELECT * FROM weather LEFT OUTER JOIN cities ON (weather.city = cities.name)
```

```
SELECT * FROM weather LEFT JOIN cities ON (weather.city = cities.name);
```

## SQL: JOIN syntax

The most general syntax (shown on the previous slide) uses the ON clause with any logical expression:

```
SELECT * FROM employees JOIN departments ON  
    (employees.dept_id = departments.dept_id);
```

If matching is performed using the *equals* operator, and matching attributes have the same names in both data sets, the USING clause may be employed:

```
SELECT * FROM employees JOIN departments USING (dept_id);
```

In addition, if *all* attributes with identical names are to be matched, we have a *natural join*:

```
SELECT * FROM employees NATURAL JOIN departments;
```

Query results may be combined using the following operators:

- UNION – appends results of second query to those of the first one,
- INTERSECT – return all rows that are in both sets,
- EXCEPT – remove rows in the second result set from the first one.

All operators remove duplicates unless used with the ALL modifier.

SQL queries may contain subqueries in various places – pay attention to the number of attributes and rows expected. Subqueries are often used with set operators:

- EXISTS – returns true if subquery has at least one row,
- IN/NOT IN – checks if value belongs/doesn't belong to set returned by subquery,
- ANY/SOME/ALL – similar to the above, but use arbitrary operators (instead of = and <>). [More info](#)



# Designing databases

---

## Good practices for database design

- Every real-world object = one entity = one relation.
- Every attribute occurs once, with its own object.
- Decompose non-atomic attributes.
- 1:n relationships: take primary key of “1”, migrate to “n” relation as foreign key.
- m:n relationships: create associative entity, use sum of primary keys from “m” and “n” as primary key.

## Implementing „inheritance”

- The notion of *subclassing* is difficult to implement in a relational database, and leads to trade-offs.
- Possible approaches:
  1. one table for superclass (and common attributes), one table for each subclass.
  2. one table for each subclass; common attributes replicated in each one. Must use UNION to obtain a set on superclass level.
  3. one table for all; each row contains attributes of all subclasses (and those of the superclass). This leads to many NULLs.

# Database normalization

- Rules to organise attributes and relations to avoid redundancy and improve integrity.
- Aimed at avoiding anomalies:
  - Update anomalies
  - Insertion anomalies
  - Deletion anomalies
- Normalisation provides a formal way to enforce good database design practices.

## Employees' Skills

Employee ID	Employee Address	Skill
426	87 Sycamore Grove	Typing
426	87 Sycamore Grove	Shorthand
519	94 Chestnut Street	Public Speaking
519	96 Walnut Avenue	Carpentry

## Faculty and Their Courses

Faculty ID	Faculty Name	Faculty Hire Date	Course Code
389	Dr. Giddens	10-Feb-1985	ENG-206
407	Dr. Saperstein	19-Apr-1999	CMP-101
407	Dr. Saperstein	19-Apr-1999	CMP-201

424	Dr. Newsome	29-Mar-2007	?
-----	-------------	-------------	---

## Faculty and Their Courses


Faculty ID	Faculty Name	Faculty Hire Date	Course Code
389	Dr. Giddens	10-Feb-1985	ENG-206
407	Dr. Saperstein	19-Apr-1999	CMP-101
407	Dr. Saperstein	19-Apr-1999	CMP-201

DELETE

# PostgreSQL


---

- PostgreSQL is a *relational database management system* (RDBMS)
- Cross-platform & open-source
- Client-server architecture:
  - server listens on 5432/tcp by default
  - clients use *libpq* to connect



# PostgreSQL concepts

- A *cluster* is one or more database instances.
- Each *database* contains at least one schema, *public*.
- Any number of schemas can be created.
- Tables and other objects are assigned to a *specific schema*.



{source: Korry Douglas, PostgreSQL, 2nd Edition}

## PostgreSQL schemas

- Creating a schema:

```
CREATE SCHEMA my_other_schema;
```

- Objects in schemas can be referred to by schema.object.
- Otherwise the DB looks in schemas listed in `search_path`, which act much like the `PATH` system environment variable.
- To set `search_path`:

```
SET search_path TO public, my_other_schema;
```

```
SET search_path TO my_other_schema, public;
```

```
SET search_path TO my_other_schema;
```



- Collection of [command-line utilities](#), including psql and several helper programs.
- Any application can use libpq, usually via some wrapper.
- GUI apps: [phpPgAdmin](#), [Adminer](#), [pgAdmin](#)

## Using *psql*: command-line switches

```
psql [option...] [dbname [username]]
```

Useful switches:

- -d dbname – specify database name
- -h hostname – specify host to connect to; if omitted, it will connect to local UNIX domain socket specified in the config
- -U username – specify username
- -l – list databases
- -c command – run specified command and exit
- -f filename – run commands from file and exit

## Using *psql*: interactive shell

Some useful meta-commands:

- `\l` – list databases
- `\d` – list objects
- `\dE`, `\di`, `\dm`, `\ds`, `\dt`, `\dv` – list objects of specific type: foreign table, index, materialized view, sequence, table, and view
- `\i` – read commands from file
- `\q` – quit

More: [psql documentation](#)